

▼ Keras -- MLPs on MNIST

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this comm
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
```

↳ Using TensorFlow backend.

```
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    plt.show()
```

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

↳ Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

```
print("Number of training examples : ", X_train.shape[0], "and each image is of shape (%d, %d)")
print("Number of training examples : ", X_test.shape[0], "and each image is of shape (%d, %d)")
```

↳ Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
# if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples : ", X_train.shape[0], "and each image is of shape (%d)"%(X_
print("Number of training examples : ", X_test.shape[0], "and each image is of shape (%d)"%(X_
```

↳ Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
# An example data point
```

```
print(x_train[0])
```

```
# if we observe the above matrix each cell is having a value between 0-255  
# before we move to apply machine learning algorithms lets try to normalize the data  
#  $X \Rightarrow (X - X_{\min}) / (X_{\max} - X_{\min}) = X / 255$ 
```

```
X_train = X_train/255  
X test = X test/255
```

```
# example data point after normalizing  
print(X_train[0])
```



```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 0, 1, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```

```
Class label of first image : 5  
After converting the output into a vector : [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
```

Softmax classifier

```
from keras.models import Sequential  
from keras.layers import Dense, Activation
```

```
# some model parameters
```

```
output_dim = 10  
input_dim = X_train.shape[1]
```

```
batch_size = 156  
nb_epoch = 20
```

```
# start building a model  
model = Sequential()
```

The model needs to know what input shape it should expect.

For this reason, the first layer in a Sequential model

```
# (for this reason, the first layer in a Sequential model  
# (and only the first, because following layers can do automatic shape inference)  
# needs to receive information about its input shape.
```

you can use input shape and input dim to pass the shape of input

```
# output_dim represent the number of nodes need in that layer  
# here we have 10 nodes
```

```
model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

Before training a model, you need to configure the learning process, which is done via the `fit` method.

```
# It receives three arguments:
```

```
# It receives three arguments:  
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io  
# A loss function. This is the objective that the model will try to minimize., https://keras.io  
# A list of metrics. For any classification problem you will want to set this to metrics=['acc']
```

```
# Note: when using the categorical_crossentropy loss, your targets should be in categorical form  
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector 1  
# for a 1 at the index corresponding to the class of the sample).
```

```
# that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_:
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss vali
# metrics values at successive epochs, as well as validation loss values and validation metric

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, vali
```



Train on 60000 samples, validate on 10000 samples
Epoch 1/20

```
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

⇒ Test score: 0.3455672636508942
Test accuracy: 0.9051

nfig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

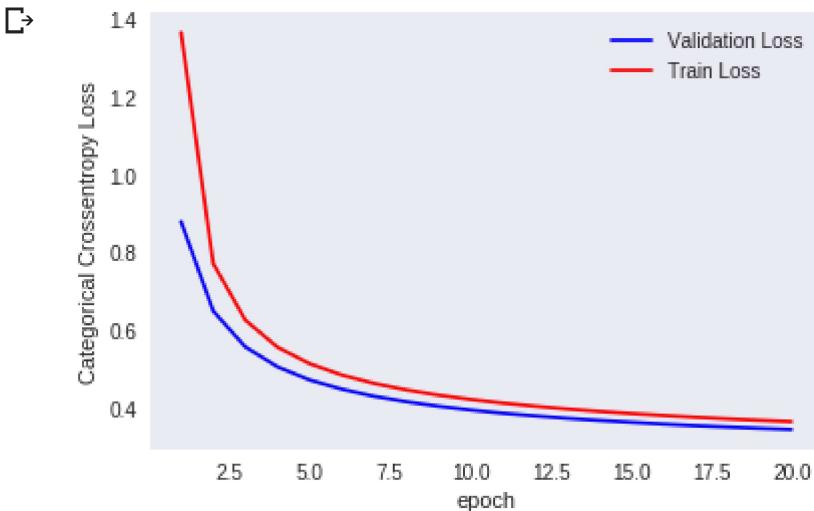
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



MLP + ReLU +Adam

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2}$ ,
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.059 \Rightarrow N(0,\sigma) = N(0,0.059)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.113 \Rightarrow N(0,\sigma) = N(0,0.113)$ 
```

```
model_relu = Sequential()
model_relu.add(Dense(556, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.1))
model_relu.add(Dense(156, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1)))
model_relu.add(Dense(output_dim, activation='softmax'))
```

model_relu.summary()

⇨

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 556)	436460
dense_3 (Dense)	(None, 156)	86892
dense_4 (Dense)	(None, 10)	1570
Total params:	524,922	
Trainable params:	524,922	
Non-trainable params:	0	

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
```

⇨

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 116us/step - loss: 0.1708 - acc: 0.9
Epoch 2/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0760 - acc: 0.9
Epoch 3/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0483 - acc: 0.9
Epoch 4/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0334 - acc: 0.9
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

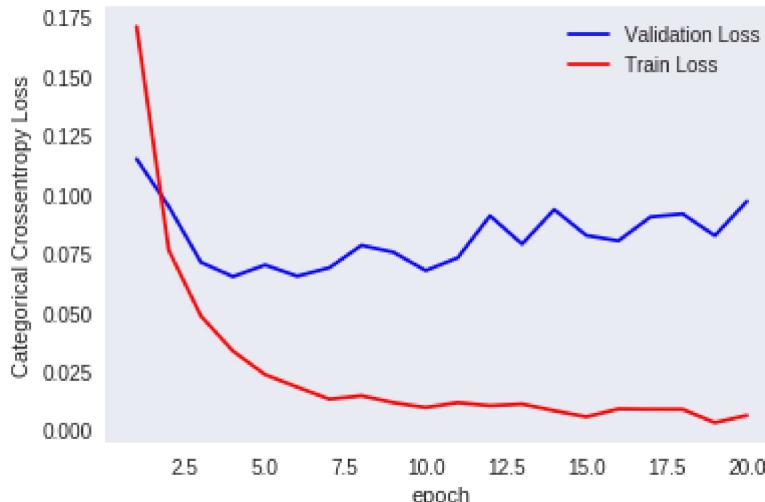
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

↳ Test score: 0.09683581481487936
 Test accuracy: 0.9802



Input(784)+ReLU(580)+ReLU(325)+ReLU(125)+Softm

```
# h1 => σ=√(2/(fan_in)) = 0.058 => N(0,σ) = N(0,0.058)
# h2 => σ=√(2/(fan_in)) = 0.078 => N(0,σ) = N(0,0.078)
# h3 => σ=√(2/(fan_in)) = 0.126 => N(0,σ) = N(0,0.126)

from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model_relu = Sequential()
model_relu.add(Dense(580, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.058)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(325, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.078)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(125, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.126)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

⇨

Layer (type)	Output Shape	Param #
<hr/>		
dense_7 (Dense)	(None, 580)	455300
<hr/>		
batch_normalization_2 (Batch Normalization)	(None, 580)	2320
<hr/>		
dropout_1 (Dropout)	(None, 580)	0
<hr/>		
dense_8 (Dense)	(None, 325)	188825
<hr/>		
batch_normalization_3 (Batch Normalization)	(None, 325)	1300
<hr/>		
dropout_2 (Dropout)	(None, 325)	0
<hr/>		
dense_9 (Dense)	(None, 125)	40750
<hr/>		
batch_normalization_4 (Batch Normalization)	(None, 125)	500
<hr/>		
dropout_3 (Dropout)	(None, 125)	0
<hr/>		
dense_10 (Dense)	(None, 10)	1260
<hr/>		
Total params: 690,255		
Trainable params: 688,195		
Non-trainable params: 2,060		

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
```

```
↳ Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 14s 229us/step - loss: 0.6143 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 13s 209us/step - loss: 0.2568 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 12s 206us/step - loss: 0.1946 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 12s 207us/step - loss: 0.1592 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 12s 207us/step - loss: 0.1401 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 13s 209us/step - loss: 0.1235 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 13s 211us/step - loss: 0.1173 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 12s 204us/step - loss: 0.1047 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 12s 207us/step - loss: 0.0983 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0920 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0849 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0810 - acc: 0.
Epoch 13/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0771 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 12s 204us/step - loss: 0.0735 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 13s 213us/step - loss: 0.0705 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0679 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0663 - acc: 0.
Epoch 18/20
60000/60000 [=====] - 12s 207us/step - loss: 0.0638 - acc: 0.
Epoch 19/20
60000/60000 [=====] - 12s 203us/step - loss: 0.0606 - acc: 0.
Epoch 20/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0578 - acc: 0.
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```

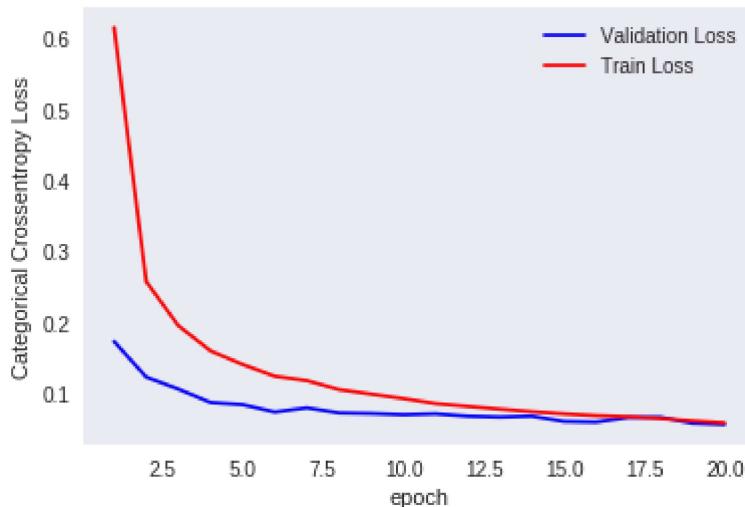
```
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

↳ Test score: 0.055547566109968464
Test accuracy: 0.9837



Input(784)+ReLU(658)+ReLU(525)+ReLU(410)+ReLU(350)+ReLU(150)+Output(10)

```
# h1 => σ=√(2/(fan_in)) = 0.055 => N(0,σ) = N(0,0.055)
# h2 => σ=√(2/(fan_in)) = 0.061 => N(0,σ) = N(0,0.061)
# h3 => σ=√(2/(fan_in)) = 0.069 => N(0,σ) = N(0,0.069)
# h4 => σ=√(2/(fan_in)) = 0.075 => N(0,σ) = N(0,0.075)
# h5 => σ=√(2/(fan_in)) = 0.115 => N(0,σ) = N(0,0.115)
```

```
model_relu = Sequential()
model_relu.add(Dense(658, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.055))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(525, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.061))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(410, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.069))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(350, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.075))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(150, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.115))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
```

```
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```



Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 658)	516530
batch_normalization_5 (Batch Normalization)	(None, 658)	2632
dropout_4 (Dropout)	(None, 658)	0
dense_12 (Dense)	(None, 525)	345975
batch_normalization_6 (Batch Normalization)	(None, 525)	2100
dropout_5 (Dropout)	(None, 525)	0
dense_13 (Dense)	(None, 410)	215660
batch_normalization_7 (Batch Normalization)	(None, 410)	1640
dropout_6 (Dropout)	(None, 410)	0
dense_14 (Dense)	(None, 350)	143850
batch_normalization_8 (Batch Normalization)	(None, 350)	1400
dropout_7 (Dropout)	(None, 350)	0
dense_15 (Dense)	(None, 150)	52650
batch_normalization_9 (Batch Normalization)	(None, 150)	600
dropout_8 (Dropout)	(None, 150)	0
dense_16 (Dense)	(None, 10)	1510
<hr/>		
Total params:	1,284,547	
Trainable params:	1,280,361	
Non-trainable params:	4,186	

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
```



```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 25s 412us/step - loss: 0.8993 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 23s 378us/step - loss: 0.3052 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 23s 379us/step - loss: 0.2283 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 23s 377us/step - loss: 0.1902 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 23s 379us/step - loss: 0.1588 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 23s 384us/step - loss: 0.1435 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 23s 381us/step - loss: 0.1310 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 23s 379us/step - loss: 0.1262 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 23s 381us/step - loss: 0.1143 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 23s 378us/step - loss: 0.1075 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 23s 378us/step - loss: 0.1022 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 23s 377us/step - loss: 0.0978 - acc: 0.
Epoch 13/20
60000/60000 [=====] - 23s 378us/step - loss: 0.0929 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 23s 379us/step - loss: 0.0907 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 23s 378us/step - loss: 0.0828 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 23s 378us/step - loss: 0.0807 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 23s 378us/step - loss: 0.0799 - acc: 0.

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.062496218884619884
Test accuracy: 0.9846

